

Report for the Project “Supervised Learning of Basis Function Coefficients for Computer-generated Speech”

Franz Papst

August 14, 2017

1 Introduction

The project “Supervised Learning of Basis Function Coefficients for Computer-generated Speech” was about trying a novel approach for speech synthesis by representing spectral parameters, used for speech synthesis, as weighted sums of basis functions. The obtained coefficients are then used for training a machine learning model which is then used for predicting the coefficients for test files. The predicted coefficients are used for recomposing the spectral parameters and generating an audio file from this recomposed spectral parameters and are plotted to have a visual representation of the values and to easily compare them with the original values.

2 Implementation of the Project

The project was implemented using the Python programming language. It offers a multitude of open libraries for scientific computing and for machine learning. In this project NumPy [1], Matplotlib [2] and scikit-learn [3] were used. For synthesising the speech signal the Speech Signal Processing Toolkit [4] was used.

As test and training data the CMU_ARCTIC databases [5] as provided by the HMM-based Speech Synthesis System Demo [6] were used.

2.1 Steps for training a Model

- Identifying all phones in the training set.
- Converting the phones into numerical values.
- Encoding all training files as basis function coefficients.
- Training the model, using the numerical phone values as \mathbf{X} matrix and the coefficients as \mathbf{y} matrix.

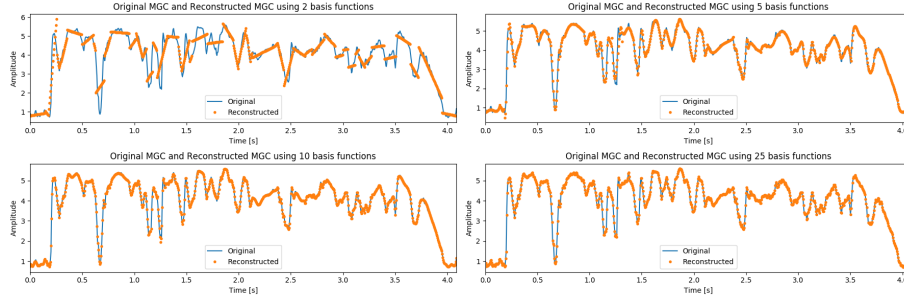


Figure 1: Different numbers of basis functions used for encoding the mel generalised cepstral coefficients (MGC).

2.2 Steps for testing a Model

- Predict the basis function coefficients using a trained model.
- Assigning the basis function coefficients to a basis function coefficients representation instance and use this instance to produce audio files and plot the predicted spectral parameters.

3 Results

3.1 Encoding using different numbers of Basis Functions

Figure 1 shows the encoding of the mel generalised cepstral coefficients (MGC) using four different (2, 5, 10 and 25) basis functions to encode the MGC values. Legendre polynomials are used for encoding the MGC values. The plots show that encoding the MGC values using only two basis functions leads to a distorted result when the values are decoded. Using 5 basis functions to encode the MGC values already leads to a result that closely resembles the original matrix. Using 10 or 25 basis functions to encode the MGC values leads to even better results, although in the produced audio files the difference is not really noticeable. Also the difference in the plots of 10 and 25 basis functions is so small that it is negligible.

Based on this findings 5 basis functions were used for encoding, because it offers a sufficient representation of the original values while not using too much resources.

3.2 Predictions of Spectral Parameters

Figures 2 and 3 show two examples of predicted MGC values done by different models. In both examples a random forest regressor gives the best results. For figure 2 the random forest regressor with 10 trees and the random forest regressor with 50 trees give almost the same prediction. For figure 3 this is true for parts

You, you would not keep the truth from me.




Figure 2: The result of different models for predicting the first line of MGC values for the phones in the sentence “You, you would not keep the truth from me.”




Figure 3: The result of different models for predicting the first line of MGC values for the phones in the sentence “Wash your hands of me.”

of the prediction, but around an x-value of 150 the random forest regressors gave different values and the one with 50 trees gives the more truthful prediction. In general the random forest regressor with 50 trees gives better results, but has also more performance requirements (especially memory). The produced audio files for regressors with 10 and 50 trees are very similar to each other and are hardly distinguishable.

The predictions by the hierarchical Gaussian mixture model (HGMM)¹ are not as good as the ones made by the random forest regressor, in both figures it the predictions made by the HGMM are vastly off the real values for some phones. In the produced audio file it is noticeable that the quality of the predictions is not as good as for the random forest regressor, audio files produced by the HGMM have the tendency to overdrive for some phones.

Both figures also show that linear regression is not really applicable for this problem. Especially at the beginning and at the end of the plots the results from linear regression are far off the real values. The prediction result is very smooth, so it is not able to capture the rapid changes of the original values (e.g. as show in figure 2 between the x-values of 300 and 400).

4 Conclusion

Spectral parameters used for resynthesis of human speech can be represented as sum of basis functions, using 5 such basis functions is enough to resynthesise a file without noticeable quality loss.

The coefficients of these basis functions were used to train machine learning models on the phones they represent in order to make predictions for given phones. Random forest regressors have shown to give the best performance compared to a hierarchical Gaussian mixture model or linear regression. A HGMM still has superior performance compared with linear regression, but the produced audio files have the tendency to overdrive.

Due to the highly non-linear nature of the problem future research on the topic could be done using neural network, especially deep neural networks.

References

- [1] E. Jones, T. Oliphant, P. Peterson *et al.*, “SciPy: Open source scientific tools for Python,” 2001–. [Online]. Available: <http://www.scipy.org/>
- [2] J. Hunter, “Matplotlib: A 2d graphics environment,” *Computing in Science Engineering*, vol. 9, no. 3, pp. 90–95, 2007. [Online]. Available: <https://matplotlib.org/>

¹A hierarchical Gaussian mixture model is a Gaussian mixture model with different layers for different combinations of phones: for a given quin-phone the model first looks up if there is the same quin-phone in the training data, if so a sample is drawn, if not the model looks up if a trained tri-phone (with the central phones of the previous quin-phone) is available, if so a sample is drawn, if not it draws a sample for the single-phone (the one which is in the centre).

- [3] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011. [Online]. Available: <http://scikit-learn.org>
- [4] S. W. Group *et al.*, “Speech signal processing toolkit (sptk),” 2006–. [Online]. Available: <http://sp-tk.sourceforge.net>
- [5] J. Kominek and A. Black, “The cmu arctic speech databases,” in *Fifth ISCA Workshop on Speech Synthesis*, 2004.
- [6] H. Zen, T. Nose, J. Yamagishi, S. Sako, T. Masuko, A. W. Black, and K. Tokuda, “The hmm-based speech synthesis system (hts) version 2.0.” in *SSW*, 2007, pp. 294–299.